

Varun Venkatesh Gowda (002126161)
Namrata Ruchandani (002125637)
Venkteshprasad Maya Rao (001087357)

Program Structures & Algorithms

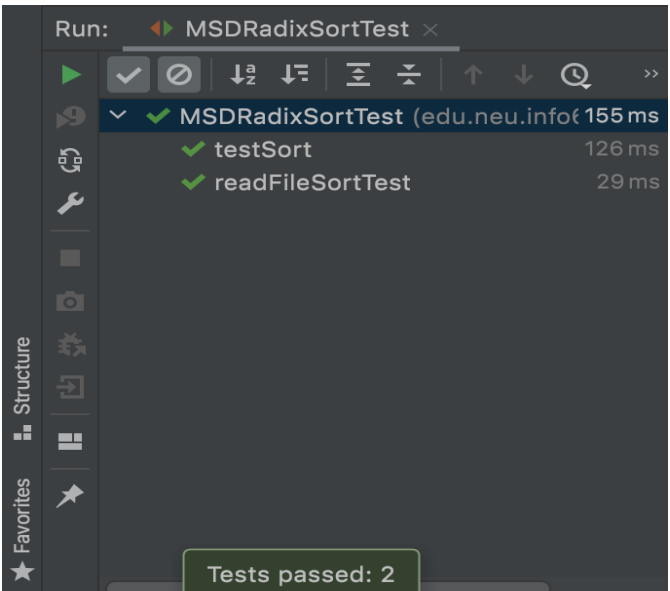

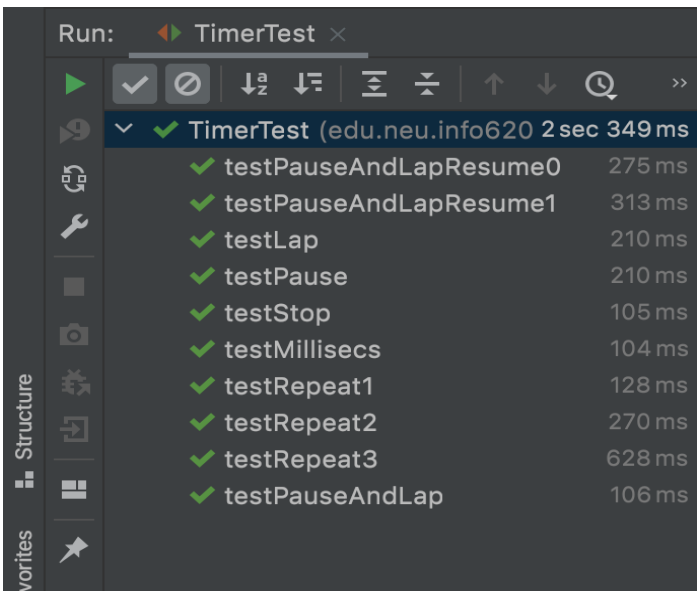
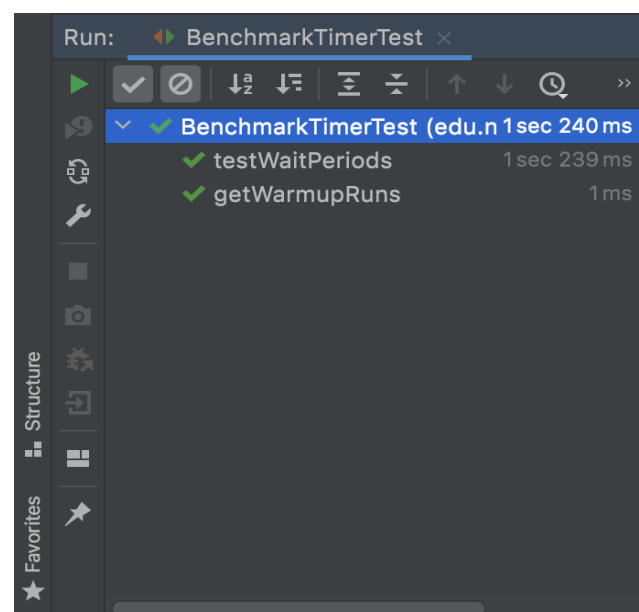
Fall 2021 Team Project

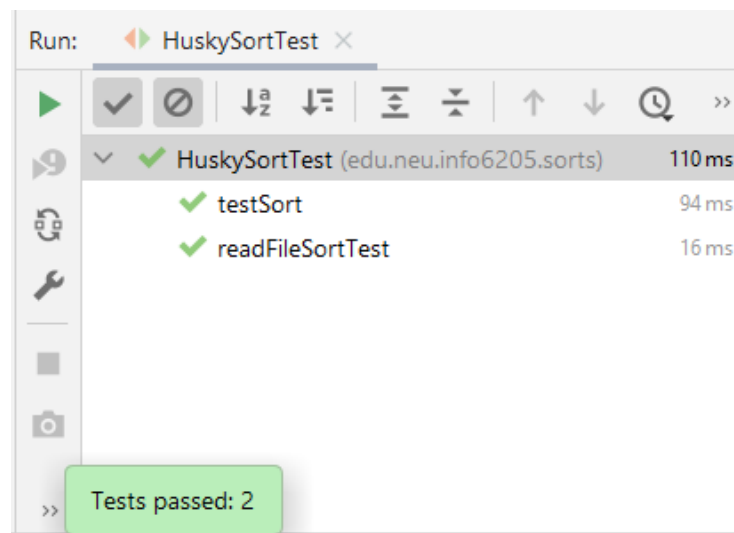
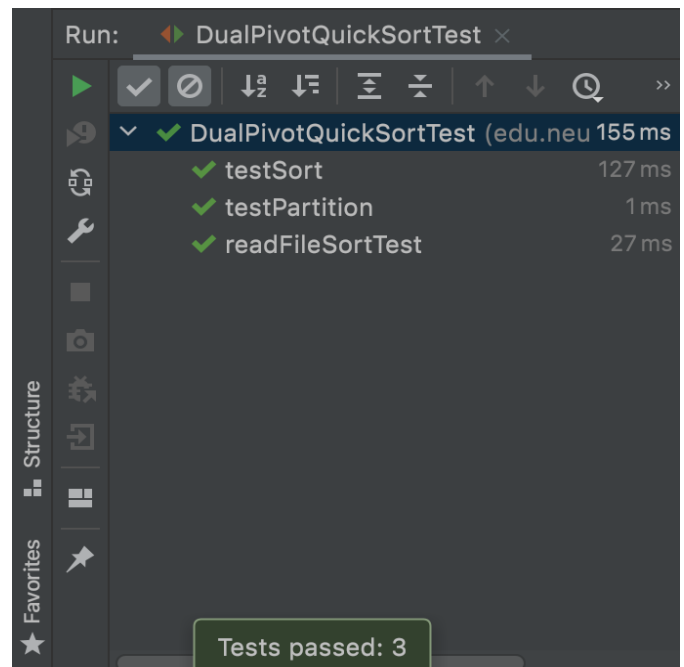
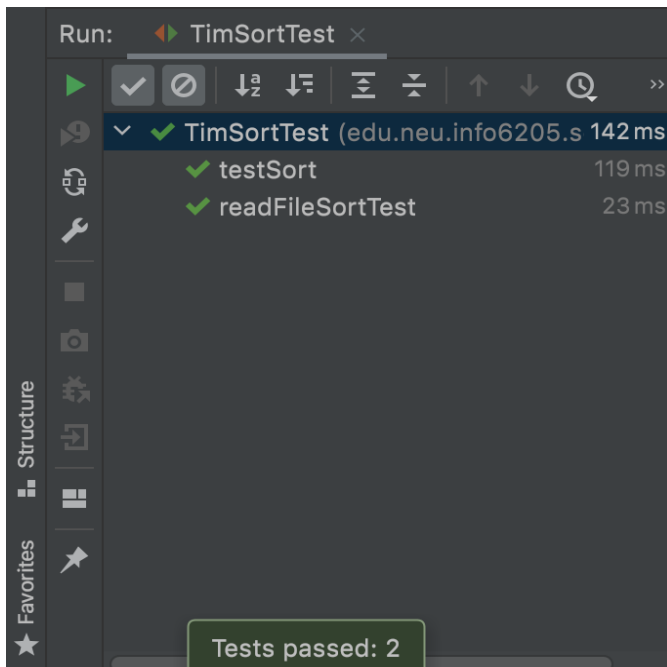
Tasks:

1. Implemented MSD Radix Sort, LSD Radix Sort, Dual Pivot Quick-Sort, Tim Sort and Husky Sort to sort an array of Chinese words in the order of the Pinyin.
2. Implemented Benchmarks in code by coding Benchmark timer to measure the running time of the sort algorithms.
3. Unit tested the sort algorithms and Benchmark utility using JUnit.
4. Implemented SortUtils for File Reader and string array generator, comparators, stubs for Chinese text and sorted Chinese text for unit tests and Chinese to pinyin converter.
5. Performed benchmarks for the sort algorithms with 250k, 500k, 1M, 2M and 4M Chinese names.
6. Recorded readings of benchmarks and tabulated the results to produce graphs for observation.
7. Prepared Report
8. Conducted readings of following papers for literature survey
 - Engineering Radix Sort by Peter M. McIlroy and Keith Bostic University of California at Berkeley; and M. Douglas McIlroy AT&T Bell Laboratories
 - PARADIS: A PARALLEL IN-PLACE RADIX SORT ALGORITHM by ROLLAND HE
 - Implementing Radixsort by Arne Anderson, Department of Computer Science, Lund University; and Stefan Nilsson, Department of Computer Science, Helsinki University of Technology
9. Examined results and prepared conclusion

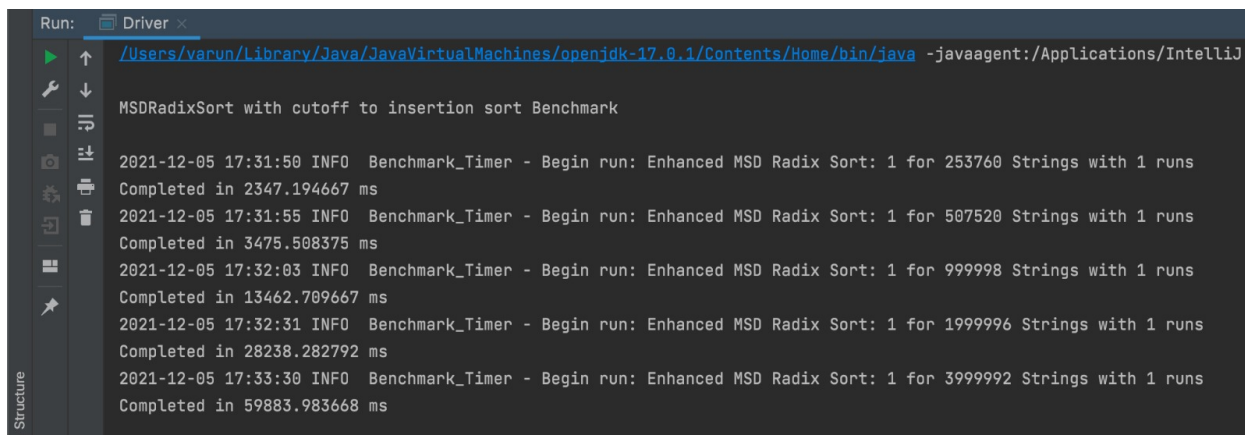
Screenshots:

Test cases passing





Benchmarking Results:



```
2021-12-05 17:33:46 INFO Benchmark_Timer - Begin run: Dual Pivot Quick Sort: 1 for 253760 Strings with 1 runs  
Completed in 785.9069 ms  
2021-12-05 17:33:49 INFO Benchmark_Timer - Begin run: Dual Pivot Quick Sort: 1 for 507520 Strings with 1 runs  
Completed in 1527.4172 ms  
2021-12-05 17:33:55 INFO Benchmark_Timer - Begin run: Dual Pivot Quick Sort: 1 for 999998 Strings with 1 runs  
Completed in 3433.5956 ms  
2021-12-05 17:34:07 INFO Benchmark_Timer - Begin run: Dual Pivot Quick Sort: 1 for 1999996 Strings with 1 runs  
Completed in 6388.9946 ms  
2021-12-05 17:34:30 INFO Benchmark_Timer - Begin run: Dual Pivot Quick Sort: 1 for 3999992 Strings with 1 runs  
Completed in 13187.5507 ms  
2021-12-05 17:35:17 INFO Benchmark_Timer - Begin run: Dual Pivot Quick Sort: 1 for 7999984 Strings with 1 runs  
Completed in 27294.7469 ms
```

Dual Pivot Quick Sort Benchmark

```
2021-12-05 15:50:49 INFO Benchmark_Timer - Begin run: Dual Pivot Quick Sort: 1 for 253760 Strings with 1 runs  
Completed in 785.9069 ms  
2021-12-05 15:50:51 INFO Benchmark_Timer - Begin run: Dual Pivot Quick Sort: 1 for 507520 Strings with 1 runs  
Completed in 1527.4172 ms  
2021-12-05 15:50:55 INFO Benchmark_Timer - Begin run: Dual Pivot Quick Sort: 1 for 999998 Strings with 1 runs  
Completed in 3433.5956 ms  
2021-12-05 15:51:04 INFO Benchmark_Timer - Begin run: Dual Pivot Quick Sort: 1 for 1999996 Strings with 1 runs  
Completed in 6388.9946 ms  
2021-12-05 15:51:19 INFO Benchmark_Timer - Begin run: Dual Pivot Quick Sort: 1 for 3999992 Strings with 1 runs  
Completed in 13187.5507 ms  
2021-12-05 15:51:51 INFO Benchmark_Timer - Begin run: Dual Pivot Quick Sort: 1 for 7999984 Strings with 1 runs  
Completed in 27294.7469 ms
```

MSD Radix Sort Benchmark

```
2021-12-05 15:52:45 INFO Benchmark_Timer - Begin run: MSD Radix Sort: 1 for 253760 Strings with 1 runs  
Completed in 1871.6247 ms  
2021-12-05 15:52:49 INFO Benchmark_Timer - Begin run: MSD Radix Sort: 1 for 507520 Strings with 1 runs  
Completed in 2770.0022 ms  
2021-12-05 15:52:55 INFO Benchmark_Timer - Begin run: MSD Radix Sort: 1 for 999998 Strings with 1 runs  
Completed in 119077.2422 ms  
2021-12-05 15:56:55 INFO Benchmark_Timer - Begin run: MSD Radix Sort: 1 for 1999996 Strings with 1 runs  
Completed in 268066.9396 ms  
2021-12-05 16:05:54 INFO Benchmark_Timer - Begin run: MSD Radix Sort: 1 for 3999992 Strings with 1 runs  
Completed in 274020.7003 ms  
2021-12-05 16:15:10 INFO Benchmark_Timer - Begin run: MSD Radix Sort: 1 for 7999984 Strings with 1 runs  
Completed in 285466.6374 ms
```

Husky Sort Benchmark

```
2021-12-05 17:33:46 INFO Benchmark_Timer - Begin run: Husky Sort: 1 for 253760 Strings with 1 runs  
Completed in 1351.4434 ms  
2021-12-05 17:33:49 INFO Benchmark_Timer - Begin run: Husky Sort: 1 for 507520 Strings with 1 runs  
Completed in 2582.9333 ms  
2021-12-05 17:33:55 INFO Benchmark_Timer - Begin run: Husky Sort: 1 for 999998 Strings with 1 runs  
Completed in 5286.9656 ms  
2021-12-05 17:34:07 INFO Benchmark_Timer - Begin run: Husky Sort: 1 for 1999996 Strings with 1 runs  
Completed in 10136.8809 ms  
2021-12-05 17:34:30 INFO Benchmark_Timer - Begin run: Husky Sort: 1 for 3999992 Strings with 1 runs  
Completed in 20350.4188 ms  
2021-12-05 17:35:17 INFO Benchmark_Timer - Begin run: Husky Sort: 1 for 7999984 Strings with 1 runs  
Completed in 44166.2954 ms
```

Process finished with exit code 0

LSD Radix Sort Benchmark

```

2021-12-05 16:24:45 INFO Benchmark_Timer - Begin run: LSD Radix Sort: 1 for 253760 Strings with 1 runs
Completed in 1088.6847 ms
2021-12-05 16:24:48 INFO Benchmark_Timer - Begin run: LSD Radix Sort: 1 for 507520 Strings with 1 runs
Completed in 1941.1791 ms
2021-12-05 16:24:52 INFO Benchmark_Timer - Begin run: LSD Radix Sort: 1 for 999998 Strings with 1 runs
Completed in 4480.2063 ms
2021-12-05 16:25:02 INFO Benchmark_Timer - Begin run: LSD Radix Sort: 1 for 1999996 Strings with 1 runs
Completed in 8430.203 ms
2021-12-05 16:25:22 INFO Benchmark_Timer - Begin run: LSD Radix Sort: 1 for 3999992 Strings with 1 runs
Completed in 17263.6091 ms
2021-12-05 16:26:02 INFO Benchmark_Timer - Begin run: LSD Radix Sort: 1 for 7999984 Strings with 1 runs
Completed in 35632.2585 ms

```

Tim Sort Benchmark




```




2021-12-05 16:27:14 INFO Benchmark_Timer - Begin run: Tim Sort: 1 for 253760 Strings with 1 runs
Completed in 771.2615 ms
2021-12-05 16:27:16 INFO Benchmark_Timer - Begin run: Tim Sort: 1 for 507520 Strings with 1 runs
Completed in 1913.6347 ms
2021-12-05 16:27:21 INFO Benchmark_Timer - Begin run: Tim Sort: 1 for 999998 Strings with 1 runs
Completed in 4128.8121 ms
2021-12-05 16:27:30 INFO Benchmark_Timer - Begin run: Tim Sort: 1 for 1999996 Strings with 1 runs
Completed in 8384.3676 ms
2021-12-05 16:27:51 INFO Benchmark_Timer - Begin run: Tim Sort: 1 for 3999992 Strings with 1 runs
Completed in 18212.8812 ms
2021-12-05 16:28:33 INFO Benchmark_Timer - Begin run: Tim Sort: 1 for 7999984 Strings with 1 runs
Completed in 36199.7933 ms

```

Benchmark times:

List Size	MSDRadixSort	LSDRadixSort	DualPivotQuickSort	Tim Sort	HuskySort	MSDRadixWithInsertion Sort
250000	1871.62	1088.68	785.9	771.26	1367.41	2347.19
500000	2770	1941.17	1527.41	1913.63	2450.6	3475.5
1000000	119077.24	4480.2	3433.59	4128.81	5437.12	13462.7
2000000	268066.93	8430.2	6388.99	8384.36	11174.77	28238.28
4000000	274020.7	17263.6	13187.55	18212.88	23806.43	59883.98
8000000	285466.63	35632.25	27294.74	36199.79	53512	





Conclusion:

The task was to implement MSD Radix Sort for comparing Chinese words and names in pinyin order, and to compare it with other sorting algorithms. In order to accomplish this, Dual Pivot Quicksort, Husky Sort, TimSort, MSD Radix Sort and LSD Radix sort were implemented.

There are instances of different Chinese characters having the same pinyin spelling (and hence the same precedence). Hence, sorting using these different algorithms produced different resultant arrays in Chinese, but their pinyin equivalents were the same. Hence, the resultant Chinese arrays were converted to pinyin to perform accuracy checks.

As per the results posted above, MSD Radix Sort performs really poorly for smaller arrays, but starts to level off as the array size reaches the 2-4 million mark. However, if we were to implement a subarray cutoff for MSD Radix sort below which insertion sort is used, performance drastically improves as shown in the screenshots.

The other sorting algorithms performed better than MSD Radix sort for the most part, with Dual Pivot Quicksort being the fastest overall. The other sorts seem to be around the same range as the DPQS, including MSD Radix with insertion sort optimization.

Credits: In addition to links provided in each file

- Prof. Robin Hillyard and contributors of the Info6205 assignments & huskysort repositories (<https://github.com/rchillyard/INFO6205>)
<https://github.com/rchillyard/The-repository-formerly-known-as>)
- pinyin4j (<http://pinyin4j.sourceforge.net/>)
- log4j (<https://logging.apache.org/log4j/2.x/>)
- ini4j (<http://ini4j.sourceforge.net/>)
- stackoverflow and geeksforgeeks community